

# Caravan EDA - MA755

Date: 02/27/2018

Author: MIMOZA MARKO (marko\_mimo@bentley.edu), SNEHA ARORA (arora\_sneh@bentley.edu), TAUSEEF AHMAD (ahmad\_taus@bentley.edu)

## Objective

The purpose of this notebook is to explain characteristics of customers. The analyses start with exploring the dataset and selecting the 10 most important variables in explaining the response variable CARAVAN. In order to perform the dimensionality reduction we apply the random forest classifier. After acquiring the important variables, we explore the single and multiple variable summaries to find patterns.

## 1. Dataset Description

The caravan insurance dataset supplied by the Dutch data mining company Sentient Machine Research and is based on a real-world business problem. The purpose of the project is to have a clear insight to why customers have a caravan policy and how these customers are different from the others.

The dataset consists of 87 attributes and 9822 observations. It is further divided into a training set (5822 observations) and a test set (4000 observations). Out of the 86 attributes, 2 are categorical, 83 are numerical and 1 is the response variable (Caravan Insurance Purchased) which indicates whether the customer purchase a caravan insurance policy or not. This is an imbalanced dataset as the target variable Caravan Insurance Purchased has more 0's i.e. the customers did not purchase the insurance policy as compared to 1's i.e. the customers did purchase the insurance policy. Furthermore, this dataset is set up as groups and each observation represent a large sample size.

There are 86 variables, containing sociodemographic data (variables 1-43) and product ownership data (variables 44-86). The sociodemographic data is derived from zip codes. All customers living in areas with the same zip code have the same sociodemographic attributes.

Each row of the dataset represents a group of customers. The following set of variables describe this group of customers:

- MAANTHUI : Number of houses 1 - 10 (in the group of customers)
- MGEMOMV : Avg size household 1 - 6 (in the group of customers)
- MOSHOOFD : Customer main type; see L2
- MOSTYPE : Customer Subtype; see L0
- MGEMLEEF : Avg age; see L1

See the documentation for the meaning of L0, L1 and L2.

<https://www.kaggle.com/uciml/caravan-insurance-challenge/data>  
(<https://www.kaggle.com/uciml/caravan-insurance-challenge/data>)

The remaining variables that start with "M" provide demographic data and contain integers between 0 and 9 that represent ranges of percentages with 0 representing 0% and 9 representing 100%. The other integers represent ranges of about 13% each. See the documentation for details.

The variables that start with "A" or "P" provide information about the customers in that postal code. They contain integers between 0 and 9 that represent ranges of counts for that variable. See the documentation for details.

The CARAVAN variable is the target variable. It is also a *count* variable (as above) recording the number of mobile home policies in that postal code.

## 2. Load Libraries

Load the required libraries and check version numbers.

```

import numpy          as np
import pandas         as pd
import matplotlib     as mpl
import matplotlib.pyplot as plt
import seaborn as sns
np.__version__, pd.__version__, mpl.__version__, sns.__version__

```

```
Out[2]: ('1.14.2', '0.18.1', '2.1.1', '0.8.1')
```

### 3. Read Dataset

First check that the files exists where we expect it.

```
%sh ls /dbfs/mnt/group-ma755/data
```

```

ais-test.csv
ais-train.csv
caravan-insurance-challenge.csv

```

Check that the file has a header and looks reasonable.

```
%sh head /dbfs/mnt/group-ma755/data/caravan-insurance-challenge.csv
```

```

ORIGIN,MOSTYPE,MAANTHUI,MGEMOMV,MGEMLEEF,MOSHOOFD,MGODRK,MGODPR,MGODOV,MGODGE,
MRELGE,MRELSA,MRELOV,MFALLEEN,MFG EKIND,MFW EKIND,MOPLHOOG,MOPLMIDD,MOPLLAAG,MBE
RHOOG,MBERZELF,MBERBOER,MBERMIDD,MBERARBG,MBERARBO,MSKA,MSKB1,MSKB2,MSKC,MSKD,
MHUUR,MHKOOP,MAUT1,MAUT2,MAUT0,MZFONDS,MZPART,MINKM30,MINK3045,MINK4575,MINK7
512,MINK123M,MINKGEM,MKOOKLA,PWAPART,PWABEDR,PWALAND,PPERSAUT,PBESAUT,PMOTSC
O,PVRAAUT,PAANHANG,PTRACTOR,PWERKT,PBROM,PLEVEN,PPERSONG,PGEZONG,PWAOREG,PBRAN
D,PZEILPL,PPLEZIER,PFIETS,PINBOED,PBYSTAND,AWAPART,AWABEDR,AWALAND,APERSAUT,AB
ESAUT,AMOTSCO,AVRAAUT,AAANHANG,ATTRACTOR,AWERKT,ABROM,ALEVEN,APERSONG,AGEZONG,A
WAOREG,ABRAND,AZEILPL,APLEZIER,AFIETS,AINBOED,ABYSTAND,CARAVAN
train,33,1,3,2,8,0,5,1,3,7,0,2,1,2,6,1,2,7,1,0,1,2,5,2,1,1,2,6,1,1,8,8,0,1,8,
1,0,4,5,0,0,4,3,0,0,0,6,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,
0,0,0,0,0,0,1,0,0,0,0,0,0,0
train,37,1,2,2,8,1,4,1,4,6,2,2,0,4,5,0,5,4,0,0,0,5,0,4,0,2,3,5,0,2,7,7,1,2,6,
3,2,0,5,2,0,5,4,2,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,
0,0,0,0,0,0,1,0,0,0,0,0,0,0
train,37,1,2,2,8,0,4,2,4,3,2,4,4,4,2,0,5,4,0,0,0,7,0,2,0,5,0,4,0,7,2,7,0,2,9,
0,4,5,0,0,0,3,4,2,0,0,6,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,
0,0,0,0,0,0,1,0,0,0,0,0,0,0
train,9,1,3,3,3,2,3,2,4,5,2,2,2,3,4,3,4,2,4,0,0,3,1,2,3,2,1,4,0,5,4,9,0,0,7,2,
1,5,3,0,0,4,4,0,0,0,6,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,0,

```

```
0,0,0,0,1,0,0,0,0,0,0
```

```
caravan_df = pd.read_csv('/dbfs/mnt/group-ma755/data/caravan-insurance-
challenge.csv')
caravan_df.head()
```

```
Out[5]:
```

	ORIGIN	MOSTYPE	MAANTHUI	MGEMOMV	MGEMLEEF	MOSHOOFD	MGODRK	MGODPR	\
0	train	33	1	3	2	8	0	5	
1	train	37	1	2	2	8	1	4	
2	train	37	1	2	2	8	0	4	
3	train	9	1	3	3	3	2	3	
4	train	40	1	4	2	10	1	4	

	MGODOV	MGODGE	...	APERSONG	AGEZONG	AWAOREG	ABRAND	AZEILPL	\
0	1	3	...	0	0	0	1	0	
1	1	4	...	0	0	0	1	0	
2	2	4	...	0	0	0	1	0	
3	2	4	...	0	0	0	1	0	
4	1	4	...	0	0	0	1	0	

	APLEZIER	AFIETS	AINBOED	ABYSTAND	CARAVAN
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	0
4	0	0	0	0	0

## Check Dataset

```
caravan_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9822 entries, 0 to 9821
Data columns (total 87 columns):
ORIGIN      9822 non-null object
MOSTYPE     9822 non-null int64
MAANTHUI    9822 non-null int64
MGEMOMV     9822 non-null int64
MGEMLEEF    9822 non-null int64
MOSHOOFD    9822 non-null int64
MGODRK      9822 non-null int64
MGODPR      9822 non-null int64
MGODOV      9822 non-null int64
MGODGE      9822 non-null int64
MRELGE      9822 non-null int64
```

```

MRELSA      9822 non-null int64
MRELOV      9822 non-null int64
MFALLEEN    9822 non-null int64
MFGEKIND    9822 non-null int64
MFW EKIND    9822 non-null int64
MOPLHOOG    9822 non-null int64
MOPLMIDD    9822 non-null int64

```

```
caravan_df.columns
```

```
Out[7]:
```

```

Index(['ORIGIN', 'MOSTYPE', 'MAANTHUI', 'MGEMOMV', 'MGEMLEEF', 'MOSHOOFD',
      'MGODRK', 'MGODPR', 'MGODOV', 'MGODGE', 'MRELGE', 'MRELSA', 'MRELOV',
      'MFALLEEN', 'MFGEKIND', 'MFW EKIND', 'MOPLHOOG', 'MOPLMIDD', 'MOPLLAAG',
      'MBERHOOG', 'MBERZELF', 'MBERBOER', 'MBERMIDD', 'MBERARBG', 'MBERARBO',
      'MSKA', 'MSKB1', 'MSKB2', 'MSKC', 'MSKD', 'MHHUUR', 'MHKOOP', 'MAUT1',
      'MAUT2', 'MAUT0', 'MZFONDS', 'MZPART', 'MINKM30', 'MINK3045',
      'MINK4575', 'MINK7512', 'MINK123M', 'MINKGEM', 'MKOOPKLA', 'PWAPART',
      'PWABEDR', 'PWALAND', 'PPERSAUT', 'PBESAUT', 'PMOTSCO', 'PVRAAUT',
      'PAANHANG', 'PTRACTOR', 'PWERKT', 'PBROM', 'PLEVEN', 'PPERSONG',
      'PGEZONG', 'PWAOREG', 'PBRAND', 'PZEILPL', 'PPLEZIER', 'PFIETS',
      'PINBOED', 'PBYSTAND', 'AWAPART', 'AWABEDR', 'AWALAND', 'APERSAUT',
      'ABESAUT', 'AMOTSCO', 'AVRAAUT', 'AAANHANG', 'ATRACTOR', 'AWERKT',
      'ABROM', 'ALEVEN', 'APERSONG', 'AGEZONG', 'AWAOREG', 'ABRAND',
      'AZEILPL', 'APLEZIER', 'AFIETS', 'AINBOED', 'ABYSTAND', 'CARAVAN'],
      dtype='object')

```

There are 87 variables and 9,822 observations in this dataframe.

```
caravan_df.shape
```

```
Out[8]: (9822, 87)
```

The 'origin' variable contains information if the observation is train or test. 5,822 of the observations are train and 4,000 are tests.

```
sum(caravan_df['ORIGIN']=='train'), sum(caravan_df['ORIGIN']=='test')
```

```
Out[9]: (5822, 4000)
```

## 4. Variable Selection

As a first step in this section we choose which variables to keep for analyzing. We first check the correlation of the numerical variables, then from the ones that are highly correlated we only keep one variable per each pair. We apply random forest classifier to check the variable importance.

```
numerical_columns = [name for name in list(caravan_df.columns) if name not in
['ORIGIN', 'CARAVAN', 'MOSTYPE', 'MOSHOOFD']]
```

Apply the `corr()` method to obtain the correlation table for the numerical columns of the `caravan_df`. The output is a pandas dataframe that we store in `correlation_df`.

```
correlation_df = caravan_df[numerical_columns].corr(method="pearson") # apply
the corr() method to numerical columns
correlation_df # pandas dataframe
```

```
Out[11]:
```

	MAANTHUI	MGEMOMV	MGEMLEEF	MGODRK	MGODPR	MGODOV	\
MAANTHUI	1.000000	-0.004315	0.052448	0.004304	-0.031296	0.020160	
MGEMOMV	-0.004315	1.000000	-0.339942	0.009383	0.053455	-0.113296	
MGEMLEEF	0.052448	-0.339942	1.000000	-0.038454	0.091775	0.058412	
MGODRK	0.004304	0.009383	-0.038454	1.000000	-0.365998	0.022942	
MGODPR	-0.031296	0.053455	0.091775	-0.365998	1.000000	-0.318426	
MGODOV	0.020160	-0.113296	0.058412	0.022942	-0.318426	1.000000	
MGODGE	0.018018	-0.010404	-0.114922	-0.079110	-0.743722	-0.130609	
MRELGE	0.006801	0.531272	-0.058155	-0.021026	0.138014	-0.134896	
MRELSA	-0.029157	-0.178395	-0.293537	0.115748	-0.215578	0.113309	
MRELOV	0.000420	-0.496816	0.212892	-0.013744	-0.068782	0.114728	
MFALLEEN	0.046156	-0.653864	0.245442	0.013623	-0.108196	0.131500	
MFGEKIND	-0.075771	-0.321324	0.212900	0.004721	0.044838	0.022794	
MFWEKIND	0.024596	0.791445	-0.371329	-0.017466	0.034227	-0.112366	
MOPLHOOG	0.005784	0.017234	-0.030063	0.244283	-0.099757	0.004198	
MOPLMIDD	-0.054934	0.039817	-0.219489	0.136335	-0.010965	0.019623	
MOPLLAAG	0.034323	-0.045710	0.176600	-0.246624	0.058878	-0.012998	
MBERHOOG	-0.004755	0.030658	0.135751	0.223981	-0.015493	-0.046798	
MBERZELF	0.034324	0.047274	0.042961	0.064208	0.039569	0.012362	
MBERBOER	-0.023410	0.102948	0.087396	-0.109850	0.115508	-0.045008	

Filter the `correlation_df` to find those variables whose absolute correlation value is greater than .7. The output is a pandas dataframe.

```
correlation_df[abs(correlation_df) > .7 ]
```

```
Out[12]:
```

	MAANTHUI	MGEMOMV	MGEMLEEF	MGODRK	MGODPR	MGODOV	MGODGE	\
MAANTHUI	1.0	NaN	NaN	NaN	NaN	NaN	NaN	
MGEMOMV	NaN	1.000000	NaN	NaN	NaN	NaN	NaN	
MGEMLEEF	NaN	NaN	1.0	NaN	NaN	NaN	NaN	
MGODRK	NaN	NaN	NaN	1.0	NaN	NaN	NaN	
MGODPR	NaN	NaN	NaN	NaN	1.000000	NaN	-0.743722	
MGODOV	NaN	NaN	NaN	NaN	NaN	1.0	NaN	
MGODGE	NaN	NaN	NaN	NaN	-0.743722	NaN	1.000000	
MRELGE	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MRELSA	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MRELOV	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MFALLEEN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MFGEKIND	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MFWEKIND	NaN	0.791445	NaN	NaN	NaN	NaN	NaN	
MOPLHOOG	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MOPLMIDD	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MOPLLAAG	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MBERHOOG	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MBERZELF	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
MBERBOER	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

Applying the `stack()` method to the to the above correlation dataframe will return all the pairs associated with the correlation magnitude. The default of the method is to remove the NaN values. Therefore, we get only the pairs with values. The output of `stack()` here is a pandas series.

```
# return pd series with the pairs of variables and their correlation magnitude
that satisfies the filter.
```

```
correlation_df[abs(correlation_df) > .7 ].stack()
```

```
Out[13]:
```

MAANTHUI	MAANTHUI	1.000000
MGEMOMV	MGEMOMV	1.000000
	MFWEKIND	0.791445
MGEMLEEF	MGEMLEEF	1.000000
MGODRK	MGODRK	1.000000
MGODPR	MGODPR	1.000000
	MGODGE	-0.743722
MGODOV	MGODOV	1.000000
MGODGE	MGODPR	-0.743722
	MGODGE	1.000000
MRELGE	MRELGE	1.000000

	MRELOV	-0.883251
MRELSA	MRELSA	1.000000
MRELOV	MRELGE	-0.883251
	MRELOV	1.000000
	MFALLEEN	0.741809
MFALLEEN	MRELOV	0.741809
	MFALLEEN	1.000000
MFGEKIND	MFGEKIND	1.000000
MFWEKIND	MGEMOMV	0.791445

The code below returns a lists of the indices with correlation magnitude higher than 0.7. The `index()` method returns the indices, which in this case are tuples. Check the above series for reference.

```
correlated_pairs=list(correlation_df[abs(correlation_df) > .7 ].stack().index)
correlated_pairs
```

```
Out[14]:
[('MAANTHUI', 'MAANTHUI'),
 ('MGEMOMV', 'MGEMOMV'),
 ('MGEMOMV', 'MFWEKIND'),
 ('MGEMLEEF', 'MGEMLEEF'),
 ('MGODRK', 'MGODRK'),
 ('MGODPR', 'MGODPR'),
 ('MGODPR', 'MGODGE'),
 ('MGODOV', 'MGODOV'),
 ('MGODGE', 'MGODPR'),
 ('MGODGE', 'MGODGE'),
 ('MRELGE', 'MRELGE'),
 ('MRELGE', 'MRELOV'),
 ('MRELSA', 'MRELSA'),
 ('MRELOV', 'MRELGE'),
 ('MRELOV', 'MRELOV'),
 ('MRELOV', 'MFALLEEN'),
 ('MFALLEEN', 'MRELOV'),
 ('MFALLEEN', 'MFALLEEN'),
 ('MFGEKIND', 'MFGEKIND'),
 ('MFWEKIND', 'MGEMOMV'),
```

Filter the `correlated_pairs` to exclude the pairs with the same variable, since it is the correlation of that variable and itself.



```
correlated_unique_pairs = [tuple for tuple in correlated_pairs if (tuple[0] !=
tuple[1])]
correlated_unique_pairs
```

```
Out[15]:
[('MGEMOMV', 'MFEWKIND'),
 ('MGODPR', 'MGODGE'),
 ('MGODGE', 'MGODPR'),
 ('MRELGE', 'MRELOV'),
 ('MRELOV', 'MRELGE'),
 ('MRELOV', 'MFALLEEN'),
 ('MFALLEEN', 'MRELOV'),
 ('MFEWKIND', 'MGEMOMV'),
 ('MOPLHOOG', 'MSKA'),
 ('MOPLMIDD', 'MOPLLAAG'),
 ('MOPLLAAG', 'MOPLMIDD'),
 ('MSKA', 'MOPLHOOG'),
 ('MHHUUR', 'MHKOOP'),
 ('MHKOOP', 'MHHUUR'),
 ('MAUT1', 'MAUT0'),
 ('MAUT0', 'MAUT1'),
 ('MZFONDS', 'MZPART'),
 ('MZPART', 'MZFONDS'),
 ('PWAPART', 'AWAPART'),
 ('PWABEDR', 'AWABEDR'),
```

The variables below has been eliminated because they were highly correlated with other variables. These variables do not provide incremental information about the response variable.

```
eliminated_variables = ['MHKOOP', 'MZFONDS', 'PWALAND', 'PWAPART', 'PGEZONG',
'PBROM', 'PBYSTAND', 'PAANHANG', 'PWAOREG', 'PZEILPL', 'PFIETS', 'PMOTSCO',
'PTRACTOR', 'PWERKT', 'PPERSAUT', 'PWABEDR', 'PVRAAUT', 'PPERSONG',
'PBESAUT', 'PPLEZIER', 'MRELOV', 'PBRAND', 'PINBOED', 'PLEVEN', 'MFEWKIND',
'MGODGE', 'MFALLEEN', 'MOPLMIDD', 'MAUT1', 'MOPLHOOG']
len(eliminated_variables)
```

```
Out[16]: 30
```

Below is displayed the list of variables that has been chosen from each pair of highly correlated variables.

```

selected_variables = ['MHHUUR', 'MZPART', 'AWALAND', 'AWAPART', 'AGEZONG',
'ABROM', 'ABYSTAND', 'AAANHANG', 'AWAOREG', 'AZEILPL', 'AFIETS', 'AMOTSCO',
'ATRACTOR', 'AWERKT', 'APERSAUT', 'AWABEDR', 'AVRAAUT', 'APERSONG',
'ABESAUT', 'APLEZIER', 'MRELGE', 'ABRAND', 'AINBOED', 'ALEVEN', 'MGEMOMV',
'MGODPR', 'MFALLEEN', 'MOPLLAAG', 'MAUT0', 'MSKA']
len(selected_variables)

```

Out[17]: 30

The final list of variables for random forest classifier includes the selected numerical variables above, the numerical variables which were not highly correlated with any other variable and the categorical variables.

```

rand_forest_var = [ var for var in caravan_df.columns if var not in
eliminated_variables and var not in ['CARAVAN', 'ORIGIN']]

```

```
len(rand_forest_var)
```

Out[19]: 55

Divide the data into train and test using only the random forest variables.

```

x_train= caravan_df.loc[caravan_df['ORIGIN']=='train', rand_forest_var]
x_test = caravan_df.loc[caravan_df['ORIGIN']=='test', rand_forest_var]

```

Define the `y_train` to be used in random forest classifier.

```
y_train=caravan_df.loc[caravan_df['ORIGIN']=='train', 'CARAVAN']
```

We perform the random forest classifier on the selected variables 55 variables.

```

from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier(max_depth=2, random_state=0) # max_depth the depth
of the tree, random_state is the random number generator
clf.fit(x_train, y_train)
RandomForestClassifier(bootstrap=True, criterion='gini') # bootstrap sampling is
done, gini because it is classification

```

Out[22]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_split=1e-07, min_samples_leaf=1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        n_estimators=10, n_jobs=1, oob_score=False, random_state=None,
                        verbose=0, warm_start=False)
```

Display the feature importance.

clf.feature\_importances\_

Out[23]:

```
array([0.          , 0.          , 0.          , 0.          , 0.02680009,
        0.          , 0.0310786 , 0.          , 0.          , 0.          ,
        0.01221705, 0.0707111 , 0.          , 0.          , 0.          ,
        0.          , 0.01423333, 0.03382929, 0.          , 0.          ,
        0.          , 0.          , 0.01973396, 0.07491502, 0.          ,
        0.          , 0.          , 0.08627368, 0.          , 0.00770941,
        0.03507074, 0.          , 0.12319854, 0.09132489, 0.08523247,
        0.          , 0.          , 0.14391677, 0.          , 0.          ,
        0.          , 0.          , 0.          , 0.          , 0.          ,
        0.06428203, 0.          , 0.          , 0.          , 0.03573934,
        0.          , 0.0437337 , 0.          , 0.          , 0.          ])
```

Zip the columns with their corresponding feature importance and store it as a list of tuples in `variable_importance`.

```
variable_importances = list(zip(x_train.columns,clf.feature_importances_))
variable_importances
```

Out[24]:

```
[('MOSTYPE', 0.0),
 ('MAANTHUI', 0.0),
 ('MGEMOMV', 0.0),
 ('MGEMLEEF', 0.0),
 ('MOSHOOFD', 0.0268000858175244),
 ('MGODRK', 0.0),
 ('MGODPR', 0.03107859638287397),
 ('MGODOV', 0.0),
 ('MRELGE', 0.0),
 ('MRELSA', 0.0),
 ('MFGEKIND', 0.012217052498740881),
 ('MOPLLAAG', 0.0707110954272083),
```

```
( 'MBERHOOG', 0.0),
( 'MBERZELF', 0.0),
( 'MBERBOER', 0.0),
( 'MBERMIDD', 0.0),
( 'MBERARBG', 0.014233331315897685),
( 'MBERARBO', 0.03382929410894868),
( 'MSKA', 0.0),
( 'MSKB1', 0.0).
```

Sort the above list according to the second item of the tuples which is the importance magnitude.

```
sorted_var_importances = sorted(variable_importances, key=lambda x: x[1],
reverse = True)
sorted_var_importances
```

```
Out[25]:
[( 'APERSAUT', 0.14391676953794413),
( 'MINKGEM', 0.12319853835130132),
( 'MKOOPKLA', 0.09132488981010675),
( 'MINKM30', 0.08627368409867325),
( 'AWAPART', 0.08523247123055566),
( 'MHHUUR', 0.07491501678054516),
( 'MOPLLAAG', 0.0707110954272083),
( 'ALEVEN', 0.06428202924369271),
( 'APLEZIER', 0.04373369830478013),
( 'ABRAND', 0.03573934000016023),
( 'MINK7512', 0.03507074026432441),
( 'MBERARBO', 0.03382929410894868),
( 'MGODPR', 0.03107859638287397),
( 'MOSHOOFD', 0.0268000858175244),
( 'MSKD', 0.01973395959074747),
( 'MBERARBG', 0.014233331315897685),
( 'MFGEKIND', 0.012217052498740881),
( 'MINK4575', 0.007709407235974831),
( 'MOSTYPE', 0.0),
( 'MAANTHUI', 0.0),
```

Top 10 most important variables according to the random forest result. The last one is the response CARAVAN.

```
most_imp_var= ['APERSAUT', 'MINKGEM', 'MKOOPKLA', 'MINKM30', 'AWAPART', 'MHHUUR',
'MOPLLAAG', 'ALEVEN', 'APLEZIER', 'ABRAND', 'CARAVAN']
```

Create the reduced dimension dataset to contain only the selected variables.

```
red_caravan_df = caravan_df[most_imp_var]
```

Display the structure of the new dataset.

```
red_caravan_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9822 entries, 0 to 9821
Data columns (total 11 columns):
APERSAUT      9822 non-null int64
MINKGEM       9822 non-null int64
MKOOPKLA      9822 non-null int64
MINKM30       9822 non-null int64
AWAPART       9822 non-null int64
MHHUUR        9822 non-null int64
MOPLLAAG      9822 non-null int64
ALEVEN        9822 non-null int64
APLEZIER      9822 non-null int64
ABRAND        9822 non-null int64
CARAVAN       9822 non-null int64
dtypes: int64(11)
memory usage: 844.2 KB
```

## 5. Explore Single Variables

### 5.1 CARAVAN - Caravan Insurance Policy

Number of mobile home policies: 0 (no policy) or 1 (having policy)

```
grouped = red_caravan_df.groupby('CARAVAN')
grouped['CARAVAN'].agg({'count': np.size,
                        'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort_values(by='percentage', ascending =False)
```

```
Out[29]:
      count  percentage
CARAVAN
```

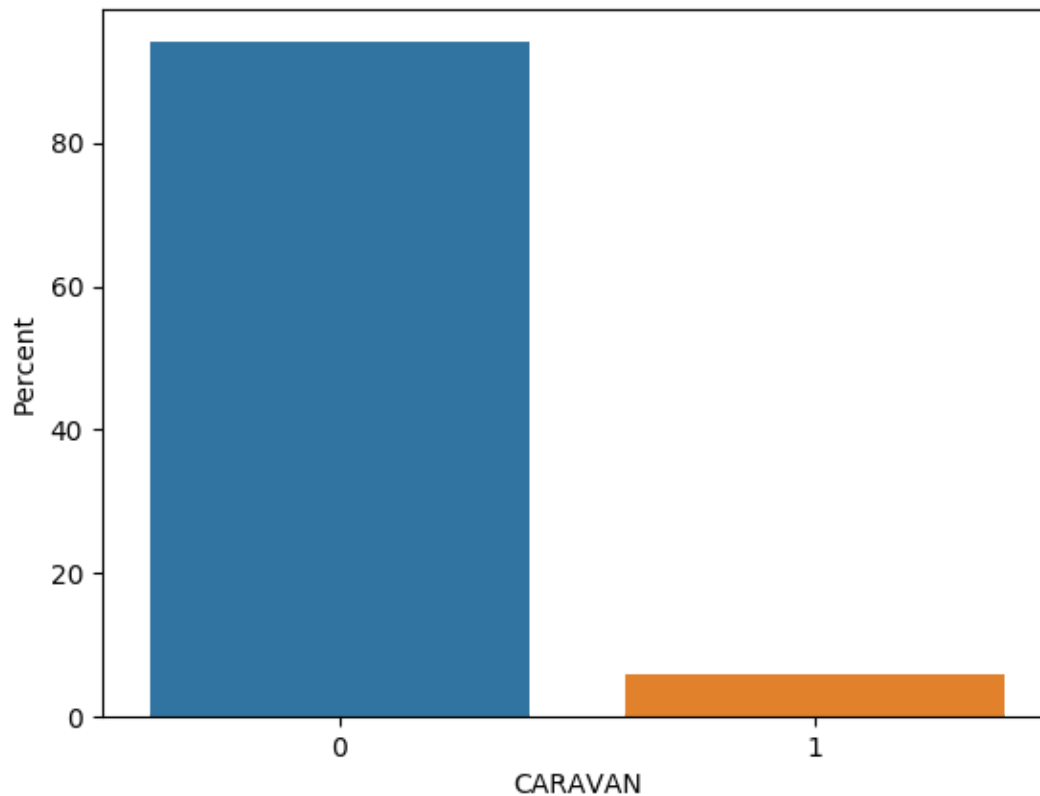
0	9236	0.940338
1	586	0.059662

6% of all the zipcodes bought caravan insurance policy.

```
caravan_plot = sns.barplot(x='CARAVAN', y='CARAVAN', data=red_caravan_df,  
estimator=lambda x: len(x) / len(red_caravan_df) * 100)  
caravan_plot.set(ylabel="Percent")
```

```
Out[30]: [Text(0,0.5,'Percent')]
```

```
display()
```



This graph visualizes the percentages of the above table. 94% of the zip codes did not purchase the insurance.

## 5.2 APERSAUT - Number of Car policies

```
grouped = red_caravan_df.groupby('APERSAUT')
grouped['APERSAUT'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort('percentage', ascending =False)
```

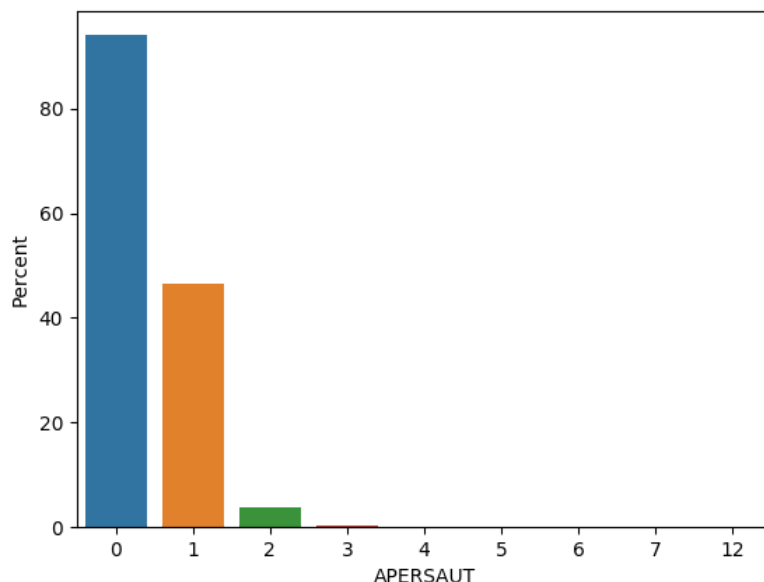
/tmp/1522871921326-0/PythonShell.py:3: FutureWarning: sort(columns=....) is deprecated, use sort\_values(by=.....)

Out[32]:

	count	percentage
APERSAUT		
0	4825	0.491244
1	4580	0.466300
2	384	0.039096
3	21	0.002138
4	8	0.000814
5	1	0.000102
6	1	0.000102
7	1	0.000102
12	1	0.000102

The above result shows that most of the zipcodes are concentrated in the low levels of car insurance. 50% of them do not have any car insurance, 47% have 1-49 car insurances and 4% have 50-99 car insurances.

```
caravan_plot = sns.barplot(x='APERSAUT', y='APERSAUT', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the above result. We can see clearly how all of the observations fall into the low levels of car insurance policies.

### 5.3 MINKGEM - Average Income

```
grouped = red_caravan_df.groupby('MINKGEM')
grouped['MINKGEM'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
                        len(red_caravan_df)}).sort('percentage', ascending =False)
```

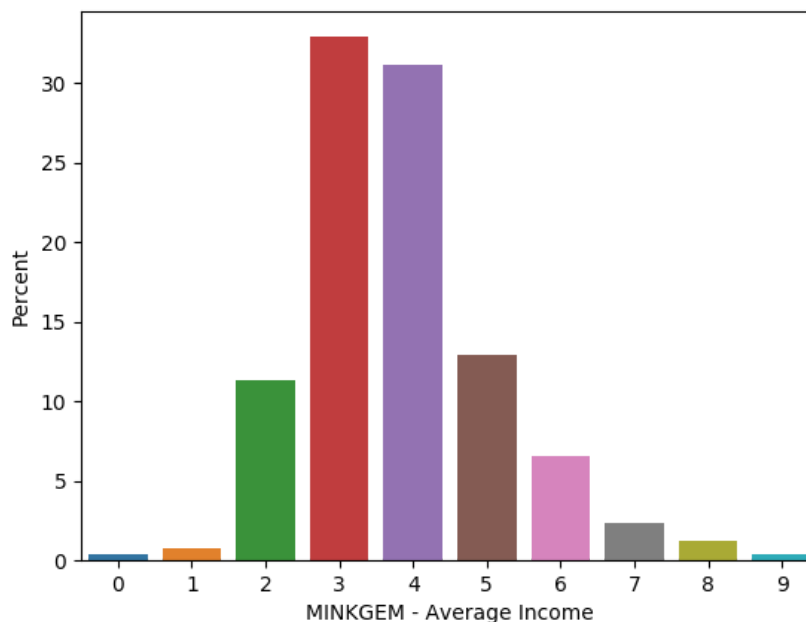
Out[34]:

	count	percentage
MINKGEM		
3	3232	0.329057
4	3063	0.311851
5	1268	0.129098
2	1110	0.113012
6	646	0.065771
7	228	0.023213
8	121	0.012319
1	78	0.007941
0	38	0.003869
9	38	0.003869



The group by result shows that 33% of the zip codes have 24-36% average income and 32% of the zip codes have 37-49% average income. Our data mostly (about 64%) contains areas with 24-49% average income.

```
caravan_plot = sns.barplot(x='MINKGEM', y='MINKGEM', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(xlabel="MINKGEM - Average Income" ,ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the average income result. We can see clearly how all of the observations fall into the middle levels average income.

## 5.4 MKOOPKLA - Purchasing Power Class

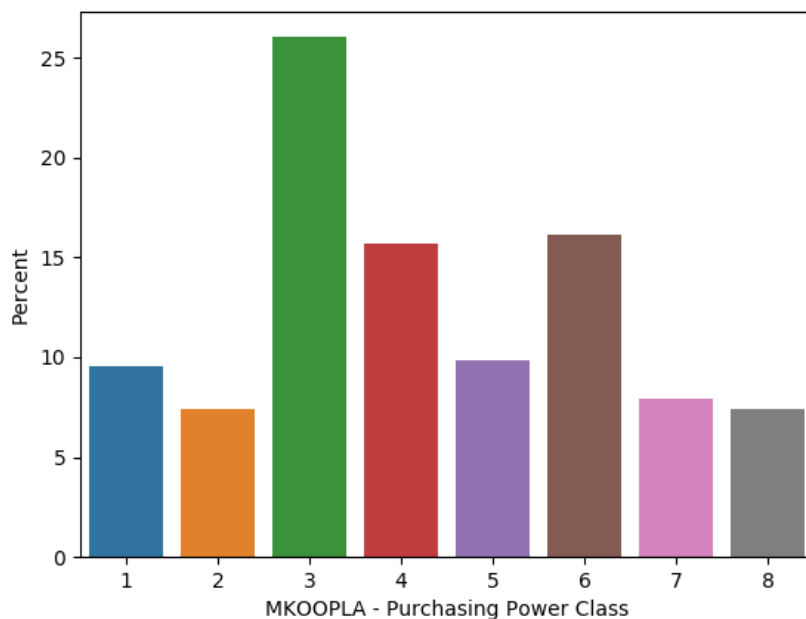
```
grouped = red_caravan_df.groupby('MKOOPKLA')
grouped['MKOOPKLA'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[107]:

	count	percentage
MK00PKLA		
3	2556	0.260232
6	1587	0.161576
4	1539	0.156689
5	964	0.098147
1	938	0.095500
7	777	0.079108
2	731	0.074425
8	730	0.074323

The group by result shows that 26% of the zip codes have 24-36% lower level education and 16% of the zip codes have 63-75% lower level education. Our data mostly contains areas with levels 3, 6 and 4 of this attribute.

```
caravan_plot = sns.barplot(x='MK00PKLA', y='MK00PKLA', data=red_caravan_df,  
estimator=lambda x: len(x) / len(red_caravan_df) * 100)  
caravan_plot.set(xlabel="MK00PLA - Purchasing Power Class", ylabel="Percent")  
display()  
plt.gcf().clear()
```



This graph visualizes the percentages of the lower level education result. We can see clearly how most of the observations fall into the 3, 6 and 4 of of this attribute.

## 5.5 MINKM30 - Income less than 30,000

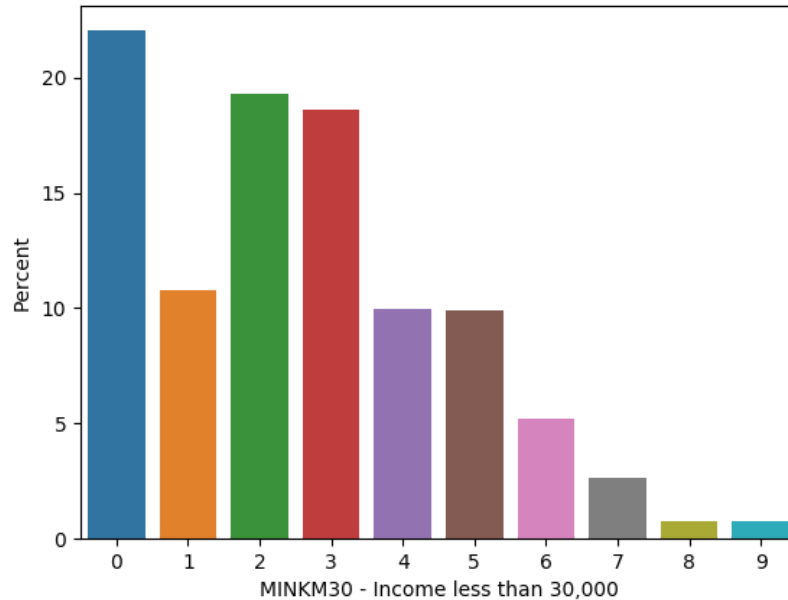
```
grouped = red_caravan_df.groupby('MINKM30')
grouped['MINKM30'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
                        len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[109]:

	count	percentage
MINKM30		
0	2164	0.220322
2	1893	0.192731
3	1826	0.185909
1	1060	0.107921
4	981	0.099878
5	971	0.098860
6	512	0.052128
7	261	0.026573
9	78	0.007941
8	76	0.007738

The group by result shows that 22% of the zip codes have 0 percentage of income less than 30,000 and 37% have 11-36% of this type of income.

```
caravan_plot = sns.barplot(x='MINKM30', y='MINKM30', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(xlabel="MINKM30 - Income less than 30,000",ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the income less than 30,000 attribute. We can see clearly how most of the observations fall into the 0, 2 and 3 of of this attribute.

## 5.6 AWAPART - Third Party Insurance

```
grouped = red_caravan_df.groupby('AWAPART')
grouped['AWAPART'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
                        len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[111]:

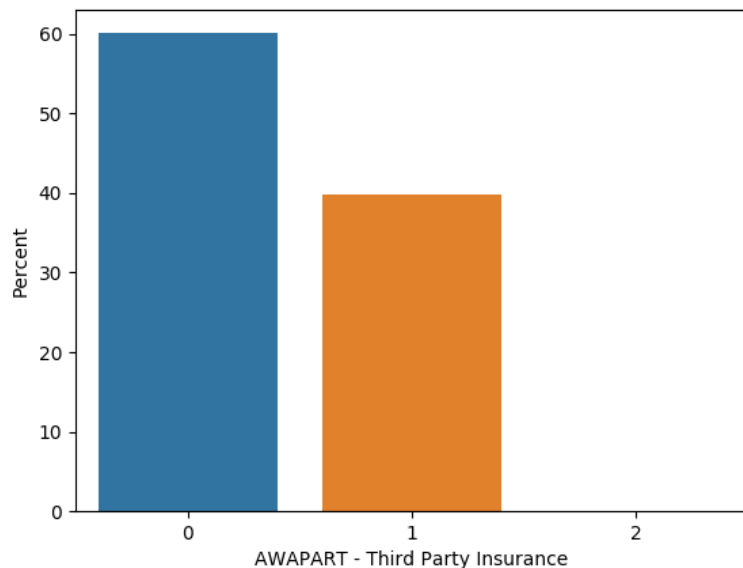
	count	percentage
AWAPART		
0	5903	0.600998
1	3909	0.397984
2	10	0.001018

The group by result shows that 60% of the zip codes have 0 third party Insurance and 39.8% of the zipcodes have between 1-49 number of third party Insurance.

```

caravan_plot = sns.barplot(x='AWAPART', y='AWAPART', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(xlabel = "AWAPART - Third Party Insurance",ylabel="Percent")
display()
plt.gcf().clear()

```



This graph visualizes the percentages of the third party Insurance attribute. We can see clearly how most of the observations most of the observations fall into the 0 and 1 level of this attribute.

## 5.7 MHHUUR - Rented House

```

grouped = red_caravan_df.groupby('MHHUUR')
grouped['MHHUUR'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort('percentage', ascending =False)

```

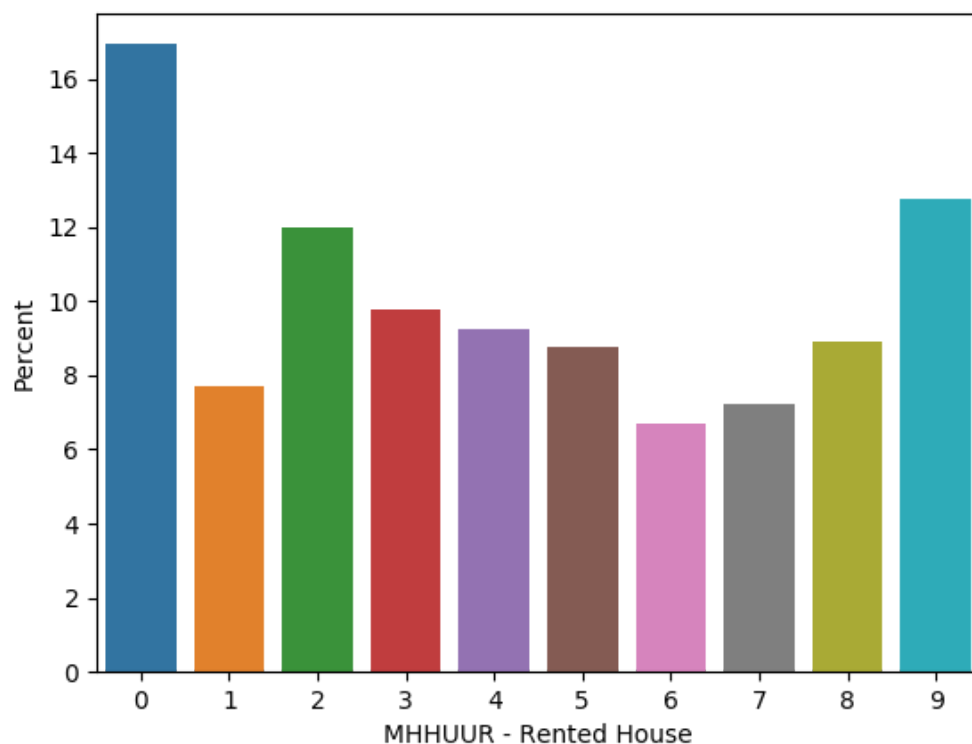
Out[113]:

	count	percentage
MHHUUR		
0	1663	0.169314
9	1255	0.127774
2	1177	0.119833
3	961	0.097842
4	908	0.092446

8	874	0.088984
5	863	0.087864
1	755	0.076868
7	710	0.072287
6	656	0.066789

The group by result shows that 17% of the zip codes have 0% of rented house and 12.7% of the zipcodes have 100% number of rented house.

```
caravan_plot = sns.barplot(x='MHHUUR', y='MHHUUR', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(xlabel="MHHUUR - Rented House",ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the rented house attribute. We can observe from the graph that 17% of the observations fall into 0 level of this attribute which means that those zipcodes has 0% of rented house and rest of the levels are almost uniformly distributed.

## 5.8 MOPLLAAG - Lower level education

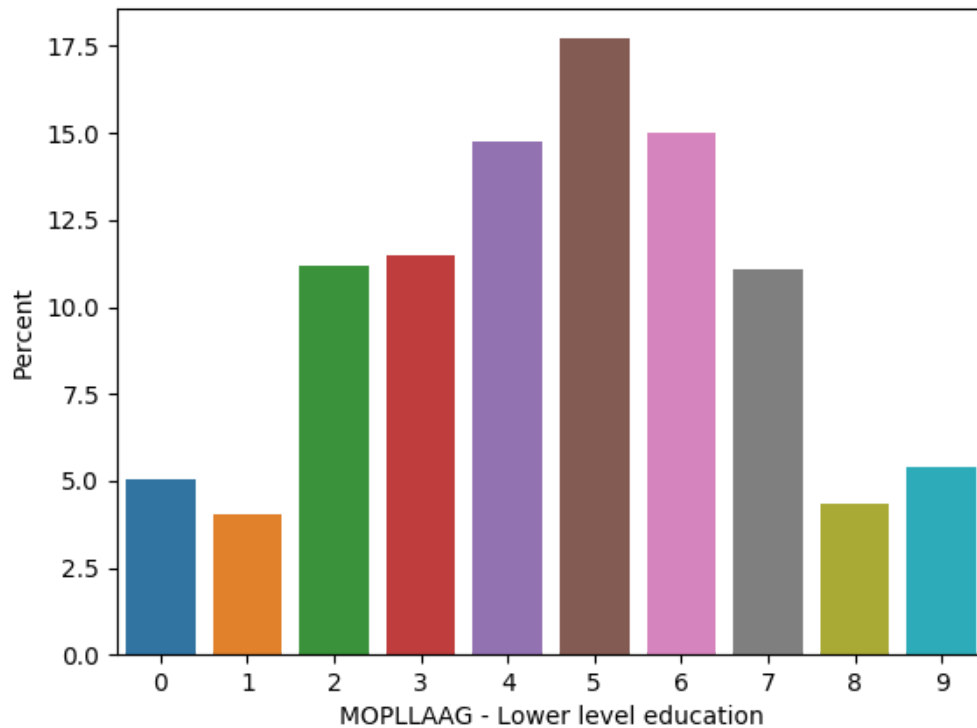
```
grouped = red_caravan_df.groupby('MOPLLAAG')
grouped['MOPLLAAG'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[115]:

	count	percentage
MOPLLAAG		
5	1740	0.177153
6	1472	0.149868
4	1448	0.147424
3	1128	0.114844
2	1099	0.111892
7	1090	0.110975
9	531	0.054062
0	494	0.050295
8	424	0.043168
1	396	0.040318

The group by result shows that 17.8% of the zip codes have 50-62% of lower level education people and 15% of the zipcodes have 63-75% number of lower level education people. We see that a high percentage of the people have lower level education in most of the zipcodes.

```
caravan_plot = sns.barplot(x='MOPLLAAG', y='MOPLLAAG', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(xlabel="MOPLLAAG - Lower level education",ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the lower level education attribute. We can observe from the graph that most of the levels for this attribute are almost uniformly distributed and most of the zipcodes have people with lower level education.

## 5.9 ALEVEN - Number of life insurances



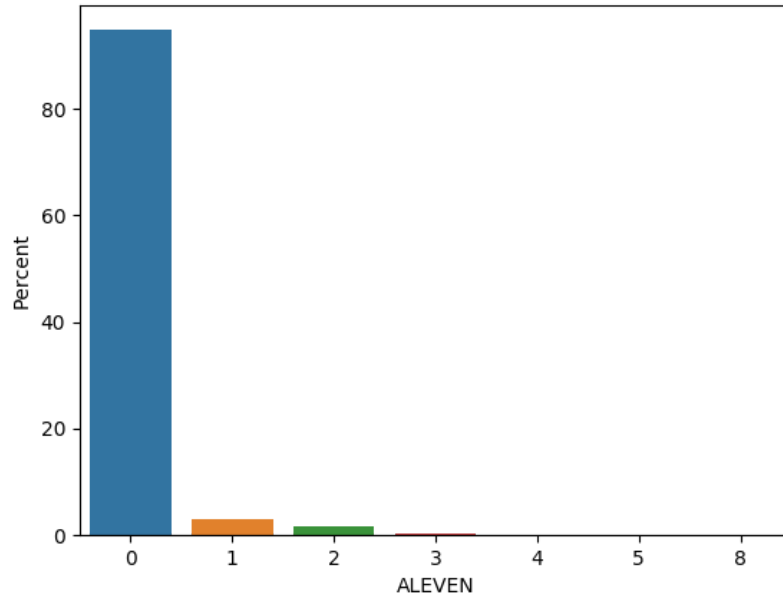
```
grouped = red_caravan_df.groupby('ALEVEN')
grouped['ALEVEN'].agg({'count':np.size,
                      'percentage': lambda x: len(x) /
                      len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[117]:

	count	percentage
ALEVEN		
0	9308	0.947668
1	305	0.031053
2	170	0.017308
3	23	0.002342
4	13	0.001324
5	2	0.000204
8	1	0.000102

The group by result shows that 95% of the zip codes dont have any Life Insurance policies and about 3.1% of the zipcodes have between 1-49 number of Life Insurance policies .

```
caravan_plot = sns.barplot(x='ALEVEN', y='ALEVEN', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the Life Insurance policy attribute. We can observe from the graph that most of the zipcodes have people with no Life insurance policy.

## 5.10 APLEZIER - Number of boat policies

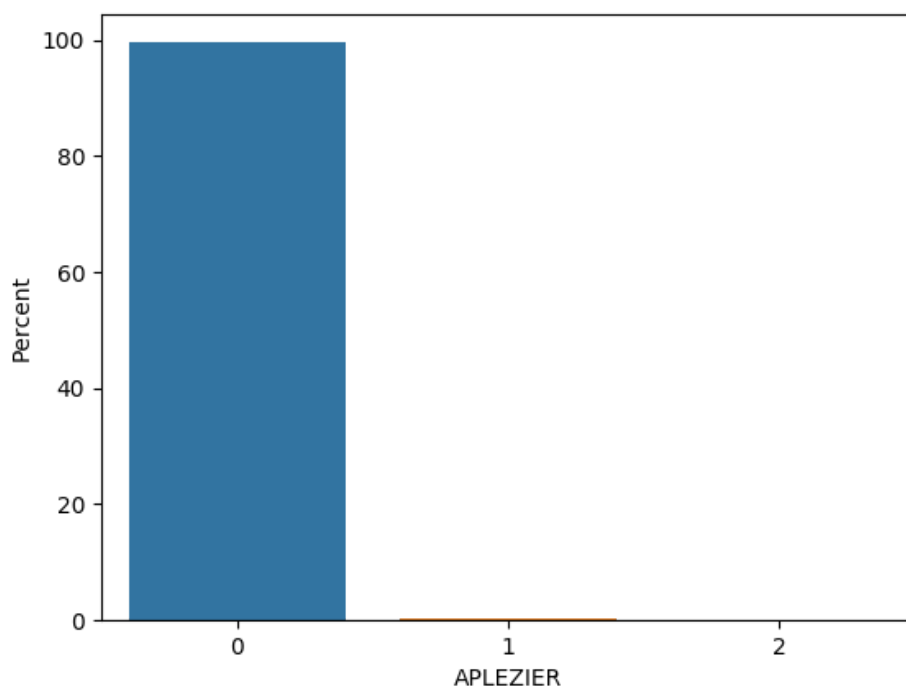
```
grouped = red_caravan_df.groupby('APPLEZIER')
grouped['APPLEZIER'].agg({'count':np.size,
                          'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[119]:

	count	percentage
APPLEZIER		
0	9777	0.995418
1	40	0.004072
2	5	0.000509

The group by result shows that 99.5% of the zip codes don't have any Boat Insurance policies and about 0.5% of the zipcodes own between 1-99 number of Boat Insurance policies .

```
caravan_plot = sns.barplot(x='APLEZIER', y='APLEZIER', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the Boat Insurance policy attribute. We can observe from the graph that most of the zipcodes have people with no boat insurance policy.

## 5.11 ABRAND - Number of fire policies

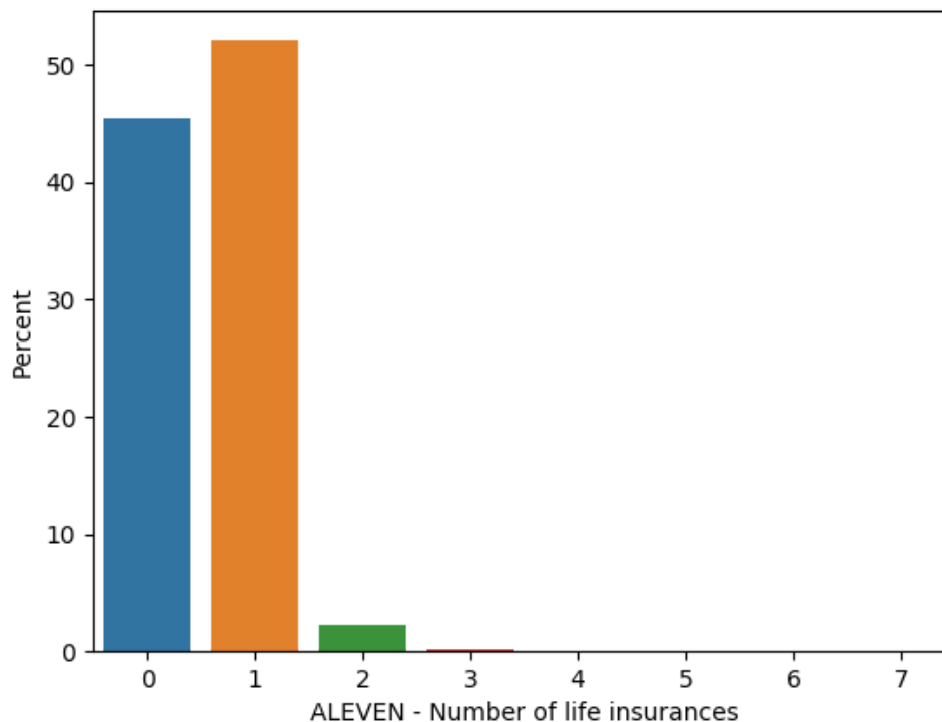
```
grouped = red_caravan_df.groupby('ABRAND')
grouped['ABRAND'].agg({'count':np.size,
                        'percentage': lambda x: len(x) /
len(red_caravan_df)}).sort('percentage', ascending =False)
```

Out[121]:

	count	percentage
ABRAND		
1	5116	0.520872
0	4464	0.454490
2	221	0.022501
3	11	0.001120
4	6	0.000611
5	2	0.000204
6	1	0.000102
7	1	0.000102

The group by result shows that 45.5% of the zip codes dont have any Fire Insurance policies and 53% of zipcodes have between 1-49 any Fire Insurance policies.

```
caravan_plot = sns.barplot(x='ABRAND', y='ABRAND', data=red_caravan_df,
estimator=lambda x: len(x) / len(red_caravan_df) * 100)
caravan_plot.set(xlabel="ALEVEN - Number of life insurances",ylabel="Percent")
display()
plt.gcf().clear()
```



This graph visualizes the percentages of the Fire Insurance policy attribute. We can observe from the graph that almost 45% of the zipcodes have people with no fire insurance policy and rest 54% of the zipcodes have 1-49 number of fire insurance policies.

## 6. Explore Multiple Variables

The purpose of the following EDA is to give a clear insight into customers having caravan insurance policy and how these customers are different from other customers.

### 6.1 Caravan vs Demographics

#### 1. Caravan, Average Income and Rented House

```
pd.pivot_table(red_caravan_df[['MHHUUR', 'CARAVAN', 'MINKGEM']],
                values = 'MINKGEM',
                index = 'MHHUUR',
                columns='CARAVAN',
                aggfunc=np.mean)
```

Out[123]:

	CARAVAN	0	1
MHHUUR			
0		4.540053	4.860606
1		4.313131	4.645161
2		4.320501	4.750000
3		3.979029	4.200000
4		3.861305	4.380000
5		3.557039	3.589744
6		3.342062	3.844444
7		3.206795	3.030303
8		2.966942	3.629630
9		3.057261	3.140000

The pivot table below shows the mean of average income for each combination of rented house and caravan insurance policy. The zip codes that bought caravan have average income between 24 to 62%. We can see that of all the zip codes that bought caravan the average income is higher for those that have lower percentage of rented house.

## 2.Caravan and Purchasing Power Class

```
grouped = red_caravan_df.groupby(['CARAVAN', 'MKOOPKLA'])
grouped['CARAVAN'].agg({'count':np.size,
                        'percentage': lambda x: len(x) / len(red_caravan_df)})
```

Out[124]:

		count	percentage
CARAVAN	MKOOPKLA		
0	1	906	0.092242
	2	712	0.072490
	3	2435	0.247913
	4	1460	0.148646
	5	909	0.092547
	6	1485	0.151191
	7	674	0.068621
	8	655	0.066687
1	1	32	0.003258
	2	19	0.001934
	3	121	0.012319
	4	79	0.008043
	5	55	0.005600
	6	102	0.010385
	7	103	0.010487
	8	75	0.007636

The result above shows all of the combinations between the caravan two levels and the purchasing power class levels in terms of count and percentages. Most zipcodes fall into caravan 0 and purchasing power class 3. We learn from these results that 25% of all the zip codes did not purchase caravan and have purchasing power class of 24-36%.

## 3.Caravan vs Lower Level Education

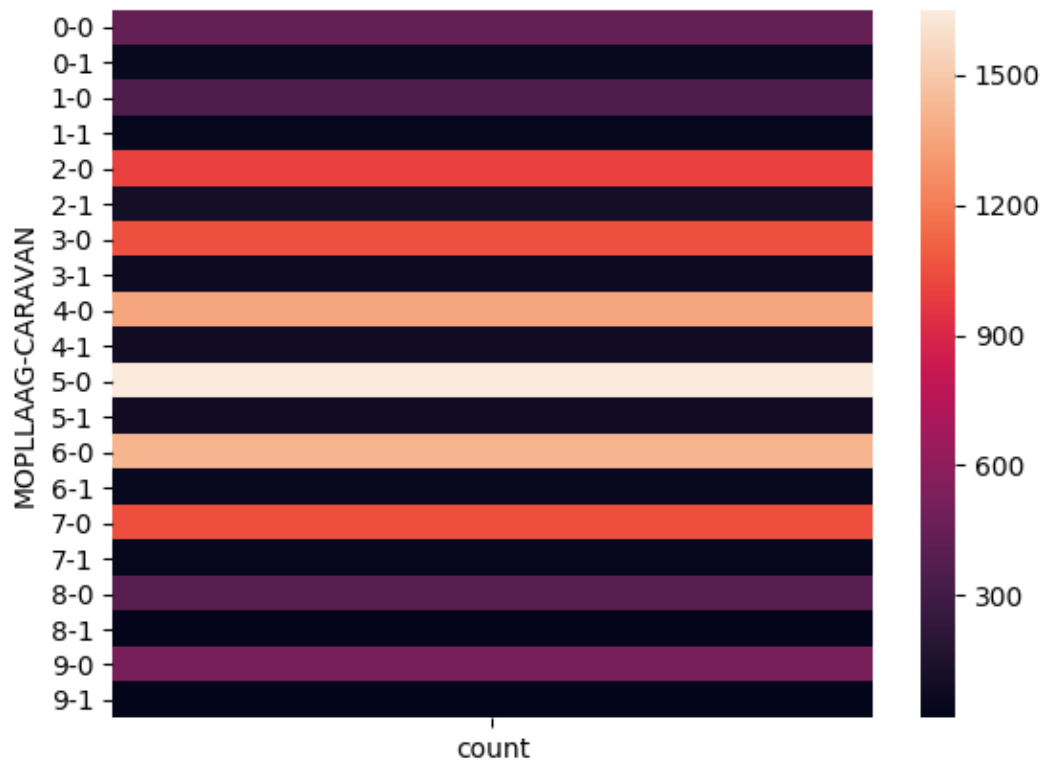
```
grouped = red_caravan_df.groupby(['CARAVAN', 'MOPLLAAG'])
grouped['CARAVAN'].agg({'count':np.size,
                        'percentage': lambda x: len(x) / len(red_caravan_df)})
```

Out[125]:

		count	percentage
CARAVAN	MOPLLAAG		
0	0	441	0.044899
	1	356	0.036245
	2	996	0.101405
	3	1055	0.107412
	4	1358	0.138261
	5	1651	0.168092
	6	1416	0.144166
	7	1051	0.107005
	8	399	0.040623
	9	513	0.052230
1	0	53	0.005396
	1	40	0.004072
	2	103	0.010487
	3	73	0.007432
	4	90	0.009163
	5	89	0.009061
	6	56	0.005701
	7	39	0.003971

The above table result shows that among the zip codes that purchased caravan policy the first comes the ones with less lower level education.

```
plt.gcf().clear()
grouped= red_caravan_df.groupby(['MOPLLAAG', 'CARAVAN'])
df=pd.DataFrame(grouped['CARAVAN'].agg({'count':np.size}))
sns.heatmap(df)
display()
plt.gcf().clear()
```



The above graph is the visualization of all the combinations of lower level education levels with those who purchased and who did not purchased caravan policy. The lighter the color the larger the number of zipcodes per that combination. The are with the most number of zipcodes correspond to the ones that did not buy caravan insurance and have lower level education of 50-62%. Also the next high number of zipcodes correspond to the ones that have lower level education of 50-75% and did not purchase caravan policy.

## 6.2 Caravan vs Other Insurances

### 1. Caravan and Car Insurance Policies



```
grouped = red_caravan_df.groupby(['CARAVAN', 'APERSAUT'])
grouped['CARAVAN'].agg({'count': np.size,
                        'percentage': lambda x: len(x) / len(red_caravan_df)})
```

Out[127]:

		count	percentage
CARAVAN	APERSAUT		
0	0	4687	0.477194
	1	4188	0.426390
	2	330	0.033598
	3	19	0.001934
	4	8	0.000814
	5	1	0.000102
	6	1	0.000102
	7	1	0.000102
	12	1	0.000102
1	0	138	0.014050
	1	392	0.039910
	2	54	0.005498
	3	2	0.000204

The result above shows all of the combinations between the caravan two levels and the car policies levels in terms of count and percentages. Most zipcodes fall into caravan 0 and car policies levels 0 and 1. We learn from these results that 90% of all the zip codes did not purchase caravan and have car policies between 0-49.

## 2.Caravan and Life insurances

```
grouped = red_caravan_df.groupby(['CARAVAN', 'ALEVEN'])
grouped['CARAVAN'].agg({'count': np.size,
                        'percentage': lambda x: len(x) / len(red_caravan_df)})
```

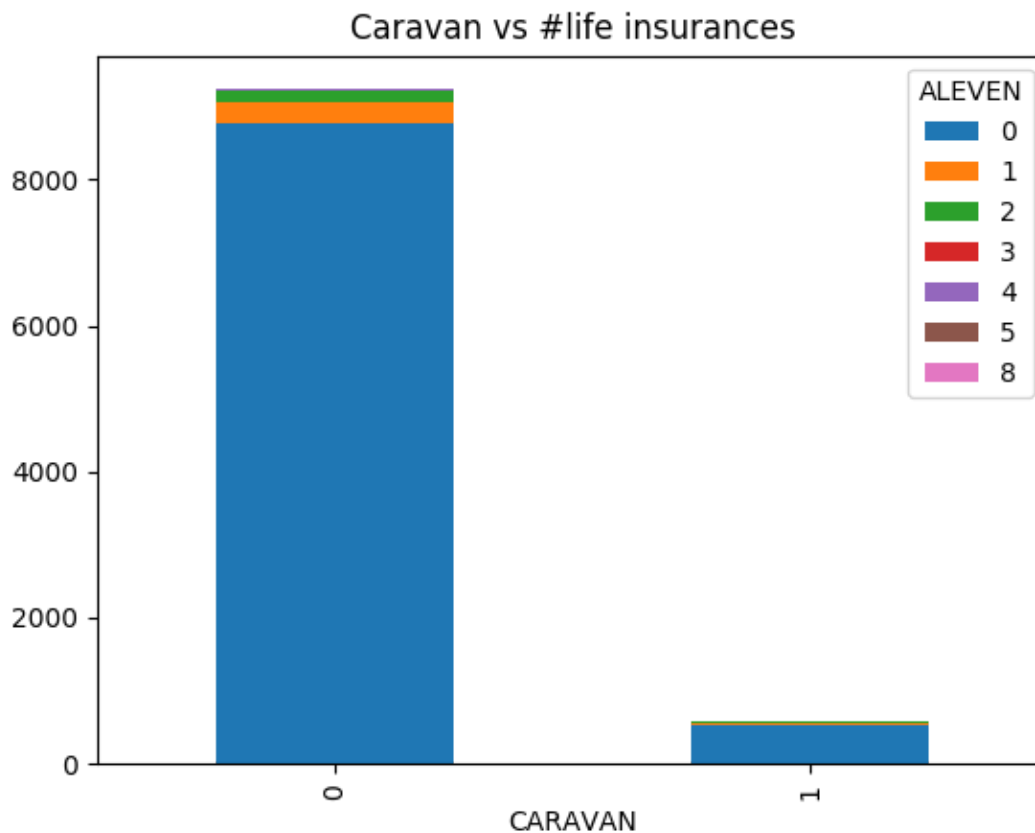
Out[128]:

		count	percentage
CARAVAN	ALEVEN		
0	0	8772	0.893097
	1	282	0.028711
	2	151	0.015374
	3	19	0.001934
	4	9	0.000916
	5	2	0.000204
	8	1	0.000102
1	0	536	0.054571

1	23	0.002342
2	19	0.001934
3	4	0.000407
4	4	0.000407

The above result shows that out of the zipcodes that fall into caravan = 0, 94% of the zipcodes do not have any life insurance policy. Similar pattern is observed for zipcodes that fall under caravan = 1.

```
df = red_caravan_df.groupby(['CARAVAN', 'ALEVEN'])
['CARAVAN'].count().unstack('ALEVEN').fillna(0)
df=pd.DataFrame(df)
df.plot(kind='bar',stacked=True,title="Caravan vs #life insurances")
display()
plt.gcf().clear()
```



The above chart shows that irrespective of the zipcode being 0 and 1, most of the zipcodes do not have any life insurance policies

### 3.Caravan and Fire Insurance Policy

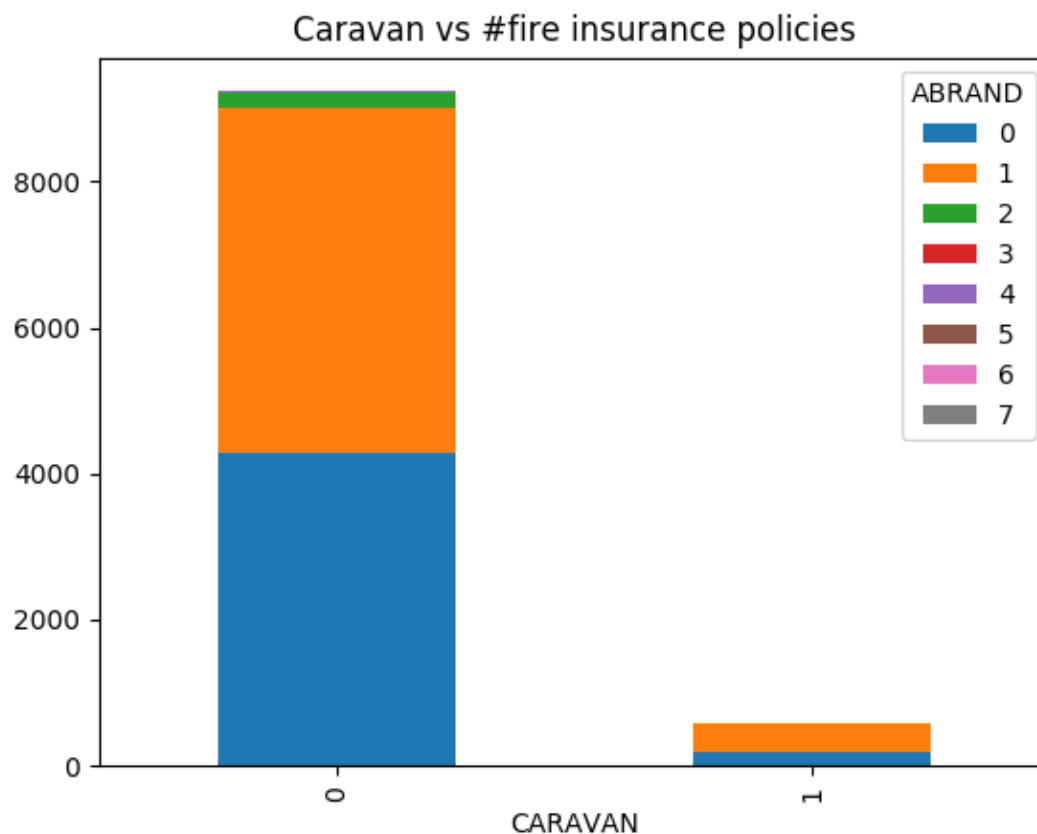
```
grouped = red_caravan_df.groupby(['CARAVAN', 'ABRAND'])
grouped['CARAVAN'].agg({'count':np.size,
                        'percentage': lambda x: len(x) / len(red_caravan_df)})
```

Out[130]:

		count	percentage
CARAVAN	ABRAND		
0	0	4283	0.436062
	1	4724	0.480961
	2	208	0.021177
	3	11	0.001120
	4	6	0.000611
	5	2	0.000204
	6	1	0.000102
	7	1	0.000102
1	0	181	0.018428
	1	392	0.039910
	2	13	0.001324

The result above shows all of the combinations between the caravan two levels and the fire policies levels in terms of count and percentages. Most zipcodes that fall into caravan 0 belong car policy levels 0 and 1. We learn from these results that 92% of all the zip codes did not purchase caravan and have fire policies between 0-49.

```
df1 = red_caravan_df.groupby(['CARAVAN', 'ABRAND'])
['CARAVAN'].count().unstack('ABRAND').fillna(0)
my_df1=pd.DataFrame(df1)
my_df1.plot(kind='bar',stacked=True,title="Caravan vs #fire insurance
policies")
display()
plt.gcf().clear()
```



The above chart shows an interesting analysis compared to the previous chart. Irrespective of whether caravan is 0 or 1, we see that there is a good proportion of zipcodes who have fire insurance.

#### 4. Caravan and Third Party Insurance

```
grouped = red_caravan_df.groupby(['CARAVAN', 'AWAPART'])
grouped['CARAVAN'].agg({'count': np.size,
                        'percentage': lambda x: len(x) / len(red_caravan_df)})
```

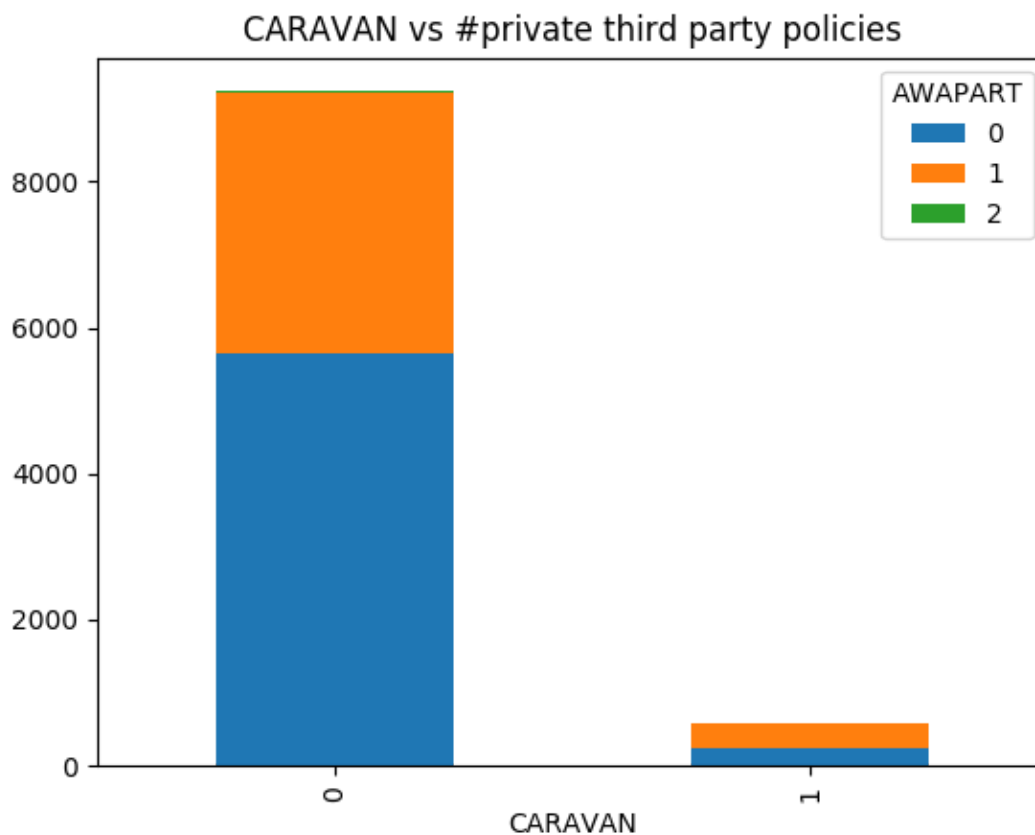
Out[132]:

CARAVAN	AWAPART	count	percentage
0	0	5656	0.575850
	1	3570	0.363470
	2	10	0.001018

1	0	247	0.025148
	1	339	0.034514

The above cross tab shows that 39% of the zipcodes, who fall under carvan = 0 category, have one or more third party insurances. While for zipcodes falling under carvan = 1, almost 58% of the them have one or more third party insurances.

```
df1 = red_caravan_df.groupby(['CARAVAN', 'AWAPART'])
['CARAVAN'].count().unstack('AWAPART').fillna(0)
my_df1=pd.DataFrame(df1)
my_df1.plot(kind='bar',stacked=True,title= "CARAVAN vs #private third party
policies")
display()
plt.gcf().clear()
```



The above chart shows similar trend as fire policies. Almost 39% of the zipcodes, who fall under carvan = 0 category have private third party insurance and 58% for those with carvan = 1.

## Conclusion

We started the analysis with the 87 variables that were available initially. We use correlation values at the first step to filter out the significant variables required for the analyses. The selected variables were uncorrelated with each other. Random forests technique was used to identify any other insignificant variables for prediction. This resulted in elimination of other insignificant variables. Some of the important factors using the random forest analyses were #fire policies, #car policies, #life insurances, #boat policies, other third-party insurances, income level and purchasing power.

Then we further drilled down to see how the final shortlisted variables vary with the dependent variable. We used bi-variate analyses including visualizations and frequency reports to see how these independent variables varied with the dependent variable.